



COPILOTTRAINING
TECHNICAL BRIEFING

The Labor Multiplier

| Where the other 67% of delivery labor lives

AI addressed the 33%. Here is where the other 67% lives — and what it costs to leave it untouched.

Executive Briefing · 30 minutes



A Three-Part Executive Briefing Series



The Labor Multiplier

What agents can actually do — the 67% of delivery labor outside the code editor that AI is built to handle at a fraction of human cost



Agentic Economics

How to calculate and capture ROI — agent compute costs, infrastructure investment requirements, and the honest payback math



Agentic Delivery

The operating model for safe scale — instrument infrastructure, pilot oversight, and the governance that enables the throughput multiplier



Each briefing is complete on its own. Together they form a full strategic framework for agentic AI investment.

Three Briefings, Three Decisions

The Labor Multiplier

Which of the 67% of delivery labor is ready to automate — and what organizational investment does capturing it require?

Opportunity Decision

Agentic Economics

What is the realistic ROI timeline, and what infrastructure makes the \$2–5/hour agent arbitrage accessible at scale?

Capital Decision

Agentic Delivery

What governance infrastructure — instruments, oversight model, no-fly zones — must exist before agent programs can operate safely?

Operating Model Decision

These are the three decisions the briefing series is designed to drive. This talk addresses one of them.

The Complete Decision Framework



Identify the Opportunity

The Labor Multiplier: 67% of delivery labor has no AI coverage — discovery, compliance, documentation, and validation are ready for agents



Quantify the Investment

Agentic Economics: agent compute at \$2–5/hour creates a structural cost arbitrage — the ROI depends on the verification infrastructure



Enable Safe Operation

Agentic Delivery: the instrument panel, pilot model, and governance structure that make the multiplier achievable at enterprise scale



The Labor Multiplier identifies what is possible. Agentic Economics quantifies the return. Agentic Delivery specifies the infrastructure that makes both achievable.

PART 1

The Hidden Labor Bill

Where AI is working — and why the gains are not reaching the P&L



PART 2

The SDLC Labor Map


Control, Coordination, Context — three buckets with a dollar figure each



PART 3

The Handoff Pattern


The mechanism that captures the gain — and why human review is non-negotiable



PART 4

The Decision

One authorization that unlocks the full operating model



The Hidden Labor Bill

AI addressed what developers type. The P&L gap lives in everything else.



The Win

84% more builds. 15% higher merge rate. The investment is real.



The Gap

Code writing is only 30–33% of engineering labor. The rest is untouched.



The Cause

AI amplifies what is already in place. The constraint has moved downstream.

Where AI investment has delivered – and where it stops

↳ Code writing: 30–33% of delivery labor

The Accenture Proof: What AI Investment Delivers

Without Copilot

Baseline build success rate

Standard PR throughput per week

Manual code review and iteration cycles

Average PR merge rate

With Copilot

84% more successful builds

8.7% more pull requests per week

Faster iteration with inline suggestions

15% higher PR merge rate

84%

more successful builds (GitHub +
Accenture 2024)

15%

higher PR merge rate across 12,000+
developers

8.7%

more pull requests per week per
developer

The Constraint Has Moved

Code generation is only the beginning

67%

of engineering labor is outside the code editor

McKinsey Global Institute, 2023

Discovery and Triage

Issue investigation, duplicate detection, impact analysis

Governance and Compliance

Security review, policy enforcement, audit generation

Knowledge Transfer

Documentation, onboarding, runbook maintenance

Planning and Coordination

Execution planning, dependency mapping, estimation

 AI addressed the 33%. The other 67% is where the P&L opportunity lives.

AI Is a Mirror and a Multiplier

Where AI Multiplies Performance

High-tempo delivery

Faster cycles get measurably faster with agents

Clean CI/CD pipelines

Stability compounds — agents amplify throughput

Accurate documentation

Agents surface and update what is already there

Where AI Reveals Gaps

Fragile build processes

Agents produce artifacts into the broken pipeline

Manual review bottlenecks

More code, same compliance review bandwidth

Documentation drift

Agents cannot synthesize what does not exist

 DORA 2025: the ROI from AI is determined more by organizational foundations than by tool selection.

What AI Investment Reaches — and What It Does Not

The AI Investment Today

Code generation and review assistance

Testing and bug-fix support in the IDE

Developer velocity in the code editor

30–33% of delivery labor addressed

The Full P&L Opportunity

Discovery, triage, and impact analysis

Compliance governance and audit generation

Knowledge transfer and documentation drift

The other 67% of delivery labor

The SDLC Labor Map

Three executive buckets — each with an addressable dollar figure.



Control

\$4.44M breach risk. Autofix: 3.7 hrs to 18 min remediation.



Coordination

18,000+ engineer-hours per year on triage. Compressed to minutes.



Context

50–200% salary to replace one engineer. 7–8 cycles per year.

Three executive buckets — each addressable today

↳ Control · Coordination · Context

Where Agent Labor Changes the Economics

CONTROL	Governance + Validation: compliance, security, and audit generation	\$4.44M avg breach
COORDINATION	Discovery + Planning: triage, impact analysis, and execution planning	18,000+ hrs/yr
CONTEXT	Knowledge + Maintenance: documentation, onboarding, team continuity	50-200% salary/hire

Control: From Manual Review to Continuous Governance

Manual Governance

4-hour compliance review per pull request

Vulnerabilities found at annual audit

Manual audit evidence assembled under deadline

Review coverage limited by developer bandwidth

Agent Governance

10-minute automated compliance check per PR

Vulnerabilities caught at the point of introduction

Audit evidence generated as a byproduct of delivery

Every PR reviewed — coverage not limited by bandwidth

12x

faster SQL injection remediation (3.7 hrs
to 18 min)

7x

faster XSS remediation (Copilot Autofix,
Octoverse 2024)

\$4.44M

average breach cost before regulatory
fines (IBM 2025)

Coordination: 18,000+ Engineer-Hours Per Year

On triage, investigation, and handoff work

18K+

engineer-hours per year on triage and coordination
— 50-person team

Atlassian State of Teams

Issue investigation

30 minutes per issue for a developer — under 5 seconds for an agent

Change impact analysis

Half a day for a senior engineer — minutes for an agent with full history

Execution planning

2–4 hours of pre-work per issue — compressed before standup

Duplicate detection

15% of issues are duplicates — agents scan the full backlog systematically



The human still makes the call — but reviews a comprehensive briefing instead of doing the research.

Context: The Compounding Cost of Forgotten Knowledge

Without Knowledge Agents

3–6 weeks to onboard a new engineer

30–60 minutes of search to understand any system

Documentation drifts from code within months

Institutional knowledge walks out with attrition

With Knowledge Agents

1–2 weeks onboarding with accurate documentation

Answers in seconds from full codebase context

Drift detected and patch proposals generated automatically

Tribal knowledge encoded and available to the entire team

💡 At 15% attrition on a 50-person team: 7–8 onboarding cycles per year — at 50–200% salary cost each (SHRM).

The Handoff Pattern

Preparation → Decision → Execution. 8–12 hours of labor. ~20 minutes of human time.



Preparation

Agent assembles context and recommendation in minutes.



Decision

Human reviews, decides, and approves. Still fully accountable.



Execution

Agent executes with monitoring. Human oversight is intact.

8–12 hours of human labor compressed
↳ ~20 minutes of human decision time

8–12 Hours to 20 Minutes: The Handoff Pattern in Practice

Without Agents

2–4 hours investigating the issue

30–60 minutes of stakeholder communication

4–8 hours of implementation work

Human time = preparation + decision + execution

With the Handoff Pattern

Agent prepares full analysis in minutes

Human reviews and decides in 2–3 minutes

Agent executes the approved plan

Human time = decision only (~20 minutes total)

~20 min

total human time per task with the
Handoff Pattern

8–12 hrs

compressed from this human-only
baseline

100%

human accountability maintained at the
decision point

Why Human Review Is Architecturally Mandatory

The trust data is the design justification

43%

of developers trust the accuracy of AI-generated output

Stack Overflow Developer Survey 2024

📊 30% trust little or none

DORA 2025: 30% of developers report little or no trust in AI outputs

🔧 The design response

Human review at the decision point is the mechanism that gives AI output integrity

⚠️ Skipping review costs more

DORA: productivity gains are lost to downstream rework when review is skipped

✅ The Handoff Pattern is the solution

Structured human review is not a bottleneck — it is what makes output trustworthy to ship

🔑 The trust gap is not a risk caveat — it is the architectural reason the human review step is non-negotiable.

DORA's 7 Foundations: The Leadership Checklist

AI Policy	Teams know what agents can do — decisions happen at the right level	Strategy
Platforms	Stable platforms — broken CI/CD means broken artifacts at agent speed	Foundation
Data Quality	Agent quality is bounded by data quality — at agent speed	Foundation
User Focus	Outcome-focused teams direct agents toward value	Culture
Team Culture	Psychological safety to flag agent errors — dysfunction gets amplified	Culture
Small Batches	Agents iterate fast; batch sizes that hide mistakes hide them longer	Delivery

The Decision

One authorization — not three asks. A governed AI policy in a bounded workflow.



The State Right Now

99% of enterprise developers already use AI tools.



The Authorization

Govern AI in one bounded high-friction workflow.



The Advantage

Governed deployments hold an 18–24 month head start.

Whether AI is governed is the remaining question
↳ 99% adoption is already the baseline

The Governance Baseline: AI Is Already in the Organization

The question is whether it is governed or unmonitored

99%

of enterprise developers have tried AI tools at work

GitHub Developer Survey 2024

The question is governance

Not whether AI is present — whether it is audited, governed, and accountable

76% using or planning AI tools

Stack Overflow 2024: adoption is already underway across the industry

The policy gap

Unmonitored adoption operates without audit trail or accountability

18–24 month advantage

Governed deployments hold a compounding head start over unstructured adoption

 The absence of a governance policy is not neutrality — it is unmonitored adoption at scale, without audit trail.

One Authorization, Three Deliverables



Labor Audit

A 4-week assessment of where engineering time actually goes — the baseline for every improvement.



90-Day Pilot

One team, three agent capabilities: triage, planning, compliance. Metrics defined before launch.



Governance Policy

Define agent autonomy boundaries, approval steps, and who is accountable in production.



The audit and pilot are implementation mechanics of one governance decision — not three separate asks.

What Changes When Agents Enter the SDLC

Before

Manual compliance review blocks or delays deployments

Triage consumes 7–8 hours of engineering time per week

Documentation drifts from code within months of writing

Developer expertise spread across investigation and execution

After

Continuous governance: every PR reviewed in under 10 minutes

Triage in seconds — humans review decisions, not raw data

Documentation synced automatically — onboarding in 1–2 weeks

Developer expertise concentrated on high-judgment decisions

12×

faster vulnerability remediation via
Copilot Autofix

~20 min

human time per 8–12 hour task with the
Handoff Pattern

\$4.44M

average breach cost addressed by
continuous governance

✓ WHAT YOU CAN DO TODAY

Today

- Identify one high-friction engineering workflow as the governance pilot target

This Week

- Commission a 4-week labor allocation assessment with VP Engineering
- Draft the governance policy scope: agent autonomy boundaries and accountability

This Month

- Launch the 90-day pilot with baseline metrics defined before day one
- Align legal and security on the AI governance policy
- Schedule a 30-day leadership review with pilot data

 Key Takeaway

Agent-assisted delivery is a current capability — authorization is the only remaining decision.

 Research and Industry Data

[McKinsey Global Institute — Developer time allocation, 2023](#)

Developers spend 30–33% of time on new code

[DORA — ROI of AI-Assisted Software Development, 2025](#)

AI as mirror and multiplier; 7 foundational capabilities

[DORA State of DevOps Report, 2023](#)

Elite teams deploy 208x more frequently; 2,600x faster recovery

[IBM Cost of a Data Breach Report, 2025](#)

\$4.44M average global breach cost before regulatory fines

 Compliance and Workforce Data

[GDPR Article 83 — Penalty Structure](#)

Up to 4% of global annual revenue for serious violations

[SHRM — Employee Replacement Cost Benchmarks](#)

50–200% of annual salary per departing engineer

 GitHub and Developer Evidence

[GitHub + Accenture Enterprise RCT, 2024](#)

84% more successful builds; 15% higher merge rate; 12,000+ developers

[GitHub Octoverse 2024](#)

SQL injection 12x faster; XSS 7x faster; 39M secrets detected

[GitHub Developer Survey 2024](#)

99% of US enterprise developers have tried AI tools at work

[Stack Overflow Developer Survey 2024](#)

43% trust AI output accuracy; 76% using or planning AI tools



The Labor Multiplier

AI ADDRESSED THE 33%. THE OTHER 67% IS THE ADDRESSABLE OPPORTUNITY.

84%

more successful builds — the AI investment is working (GitHub + Accenture 2024)

67%

of engineering labor untouched — the full P&L opportunity

~20 min

of human time per task with the Handoff Pattern in place

What is one high-friction workflow where the Handoff Pattern could operate in the next 90 days?