



COPILOTTRAINING  
TECHNICAL BRIEFING

# The Agentic Operating Model

| **How leaders bound agent authority and prove the return**

One controlled pilot. One decision record. Scale once, adjust once, or stop.

---

Executive Briefing · 30–45 minutes



## "Will leadership authorize one controlled operating-model pilot with a falsifiable return gate?"

The charter commits one team, one workflow, local funding, control capacity, and bounded agent authority.

**The return decision is recorded in advance: scale once, adjust once, or stop.**

### Authorizing Body

Approves the completed charter, funding source, and calendar return date.



### Operational Owner

The VP of Engineering owns workflow execution and measured results.



### Measurement Custody

Finance, Strategy, or an independent data owner protects baseline integrity.



### 4 weeks

Internal baseline before assisted execution begins.

### 90 days

Bounded pilot with predeclared thresholds and guardrail floors.

### 3 outcomes


Scale once, adjust once, or stop at the dated return gate.

PART 1

## The Decision

Authorize the test

Define the owner, scope, authority boundary, and local commitments.



PART 2

## The Evidence

Earn the evidence

Use external variance to design a credible local test.




PART 3

## Guardrails and Economics

Bound the operating model

Assign retained authority, controls, and fully burdened economics.



PART 4

## The Gate

Make the decision falsifiable

Predeclare thresholds, custody, stop conditions, and the return date.



# The Decision

Leadership pre-authorizes a controlled decision process with named ownership and local commitments.



## Authorization

One completed charter and accountable signatory.



## Scope

One named team and one repeatable workflow.



## Return

A dated gate aligned to a real planning cycle.

The authorization commits capacity and bounded permissions  
↳ 4-week baseline → 90-day pilot → recorded decision

## One Controlled Decision Process



### Executive Authority

A named body approves the charter, local ceiling, funding source, and tie-break authority.



### Operational Ownership

The CTO sponsors. The VP of Engineering owns the selected workflow and its results.



### Independent Custody

Finance, Strategy, or an independent data owner protects baseline and gate integrity.



The authorization buys one bounded test and one recorded return decision.



## The Charter Records the Commitments

|                   |                                           |                 |
|-------------------|-------------------------------------------|-----------------|
| <b>Authorizer</b> | Named committee and accountable signatory | <b>Decision</b> |
|-------------------|-------------------------------------------|-----------------|

|                |                                                      |                  |
|----------------|------------------------------------------------------|------------------|
| <b>Sponsor</b> | CTO owns cross-functional sponsorship and escalation | <b>Executive</b> |
|----------------|------------------------------------------------------|------------------|

|              |                                                    |                  |
|--------------|----------------------------------------------------|------------------|
| <b>Owner</b> | VP Engineering owns workflow execution and results | <b>Operating</b> |
|--------------|----------------------------------------------------|------------------|

|              |                                            |                |
|--------------|--------------------------------------------|----------------|
| <b>Scope</b> | One named team and one repeatable workflow | <b>Bounded</b> |
|--------------|--------------------------------------------|----------------|

|               |                                                    |              |
|---------------|----------------------------------------------------|--------------|
| <b>Budget</b> | Local ceiling, funding source, and labor treatment | <b>Local</b> |
|---------------|----------------------------------------------------|--------------|

|               |                                                   |              |
|---------------|---------------------------------------------------|--------------|
| <b>Return</b> | Exact date aligned to a planning or funding cycle | <b>Dated</b> |
|---------------|---------------------------------------------------|--------------|

*Control capacity, agent permissions, retained authority, and measurement custody are completed before launch.*

# The Evidence

External variance creates a positive case for a bounded local test with independent measurement.



## Positive Outcomes

Enterprise results show meaningful potential.



## Observed Variance

An expert cohort recorded slower completion.



## Local Proof

A comparable baseline determines the decision.

Current studies examine assistant-era interventions  
↳ Local agent outcomes require a local test

## Current Research Shows Outcome Variance

### GitHub and Accenture

#### **84% more successful builds**

Reported among Copilot users in an enterprise study

#### **8.7% more pull requests**

Reported per developer in the participating population

#### **15% higher merge rate**

Assistant-era result with commercial interest disclosed

### METR

#### **19% longer completion**

Randomized study of eligible tasks in mature repositories

#### **16 experienced developers**

Independent, small expert open-source cohort

#### **Expected faster completion**

Participant expectations differed from measured outcomes

 Both studies examine assistant-era human authorship. Bounded agent execution remains untested; outcomes vary by task, context, and intervention.



## Four Evidence Classes Protect the Decision



### External Observation

Calibrates plausible outcomes and limits from cited research.



### Internal Baseline

Records local workflow performance before assisted execution.



### Modeled Scenario

Applies disclosed local assumptions to fully burdened cost and value.



### Proposed Threshold

Defines the approved test for scale, adjust, or stop.

## Three Zones Bound the Candidate Work

**Control**

Checks, evidence collection, and policy-constrained preparation

**Risk****Coordination**

Triage, dependency mapping, status, and release preparation

**Flow****Context**

Repository analysis, documentation, and decision history

**Judgment**

*Enterprise ROI, universal productivity, headcount reduction, and autonomous production authority remain outside the pilot conclusions.*

# Guardrails and Economics

Bounded permissions, retained human authority, and fully burdened local economics define the test.



## Least Privilege

Access is limited to the named workflow.



## Human Authority

Accountable leaders retain reserved actions.



## Local Economics

The model includes control, review, and rework.

**Controls produce evidence throughout execution**  
↳ Any authority or control breach → immediate stop

## Authority Is Explicit and Enforceable

### Agent Scope

#### **Gather context**

Read approved repositories, records, and decision history

#### **Prepare work**

Propose plans and modify approved development assets

#### **Produce evidence**

Run approved checks and assemble review records

### Retained Human Authority

#### **Production**

Deployment, data changes, configuration, and integrations

#### **Access and controls**

Permission changes, exceptions, and new dependencies

#### **Accountability**

Objectives, risk acceptance, release, and boundary changes



Tool permissions, delivery gates, logs, and review records make the boundary observable.

## The Economic Case Is Local and Fully Burdened

### Local Case

**Measured volume**

Workflow events and validated outcomes from the baseline

**Full cost**

Tools, platform, controls, review, rework, and owner time

**Defensible value**

Redeployed capacity, flow, and quality with named destinations

### Downside Case

**Lower output**

Fewer validated outcomes than the local case

**Higher burden**

More review, rework, controls, and operating cost

**Limited redeployment**

Capacity value is reduced or unrealized

 Every assumption carries an owner, source, date, and sensitivity range.



# The Gate

A dated threshold memo makes the day-90 decision falsifiable before assisted execution begins.



## Thresholds

Outcome, economics, and guardrail floors.



## Custody

Independent ownership protects the record.



## Disposition

Scale once, adjust once, or stop.

The memo is approved after baseline and before launch  
↳ Every threshold is complete and dated

## The Threshold Memo Makes the Test Falsifiable

|                |                                                            |               |
|----------------|------------------------------------------------------------|---------------|
| <b>Minimum</b> | Completed-outcome count supported by expected event volume | <b>Volume</b> |
|----------------|------------------------------------------------------------|---------------|

|                   |                                                        |               |
|-------------------|--------------------------------------------------------|---------------|
| <b>Comparable</b> | Predefined task rule or a concurrent comparison cohort | <b>Design</b> |
|-------------------|--------------------------------------------------------|---------------|

|                |                                                    |                |
|----------------|----------------------------------------------------|----------------|
| <b>Outcome</b> | One workflow-specific result with a business owner | <b>Primary</b> |
|----------------|----------------------------------------------------|----------------|

|                  |                                                       |              |
|------------------|-------------------------------------------------------|--------------|
| <b>Economics</b> | Fully burdened threshold using disclosed local inputs | <b>Value</b> |
|------------------|-------------------------------------------------------|--------------|

|                   |                                                  |                |
|-------------------|--------------------------------------------------|----------------|
| <b>Guardrails</b> | Quality, control, authority, and evidence floors | <b>Control</b> |
|-------------------|--------------------------------------------------|----------------|

|                |                                                    |                 |
|----------------|----------------------------------------------------|-----------------|
| <b>Arbiter</b> | Named authority for a tied or disputed disposition | <b>Decision</b> |
|----------------|----------------------------------------------------|-----------------|

*The memo also records the directional floor, exact return date, and independent measurement custodian.*

## Three Dispositions Close the Decision



### Scale Once

Primary outcome, economics, volume, comparability, and every guardrail floor are satisfied.



### Adjust Once

The directional floor holds, controls remain intact, and one bounded correction is funded and dated.



### Stop

A stop condition occurs, a floor fails, comparison breaks, or the adjustment misses its gate.

● Any authority-boundary or control breach stops the pilot immediately.

## The Final Request Is One Recorded Authorization



### Complete

Fill every local charter and threshold field after the four-week baseline.



### Authorize

Approve the named team, workflow, budget, controls, permissions, and 90-day pilot.



### Return

Reconvene on the recorded date for scale once, adjust once, or stop.



The decision record preserves evidence classes, scenario assumptions, control ownership, and disposition.

## Three Actions Prepare the Authorization



### Name the Test

Select one team, one repeatable workflow, the VP Engineering owner, and the executive sponsor.



### Commission the Baseline

Assign independent measurement custody and collect four weeks of comparable workflow evidence.



### Complete the Memo

Set local budget, capacity, permissions, thresholds, guardrail floors, arbiter, and return date.



Launch follows a complete charter and dated threshold memo with every local field approved.



 **Outcome Evidence**

[GitHub and Accenture enterprise study, 2024](#)

Commercially interested assistant-era build and pull-request outcomes.

[METR experienced open-source developer study, 2025](#)

Independent randomized study reporting slower completion in a small expert cohort.

 **Governance**

[NIST AI Risk Management Framework, 2023](#)

Governance vocabulary for oversight, monitoring, and accountability.

 **Measurement Design**

[DORA ROI of AI-Assisted Software Development, 2025](#)

Commercially interested framework for outcomes and organizational conditions.

[DX engineering AI measurement practices, 2024](#)

Commercially interested review of outcome, quality, and experience measures.



# The Agentic Operating Model

ONE CONTROLLED PILOT, ONE DECISION RECORD.

## 4 weeks

Internal baseline with independent measurement custody

## 90 days

One bounded operating-model pilot

## 3 outcomes

Scale once, adjust once, or stop

Which named workflow has enough volume, control capacity, and measurable value for this authorization?